

Practical 3

Data Cleaning and Preprocessing

Aim: To clean and preprocess a dataset by handling missing values, removing duplicates, and normalizing data.

Steps:

- 1. Import a dataset (e.g., a CSV file).**
- 2. Inspect the dataset for missing values and duplicates.**
- 3. Handle missing values by filling them with mean/median or dropping rows.**
- 4. Normalize a numerical column to a scale of 0 to 1.**
- 5. Display the cleaned dataset.**

Code:

```
import pandas as pd
from sklearn.preprocessing import MinMaxScaler

# Step 1: Load dataset
df = pd.read_csv("sample_dataset.csv")

# Step 2: Inspect dataset
print(df.info())

# Step 3: Handle missing values in numeric columns
numeric_cols = df.select_dtypes(include=['number']).columns # Get numeric columns
df[numeric_cols] = df[numeric_cols].fillna(df[numeric_cols].mean()) # Fill NaN with column mean

# Step 4: Remove duplicates
df.drop_duplicates(inplace=True)

# Step 5: Normalize a specific numerical column
scaler = MinMaxScaler()
df['Normalized_Column'] = scaler.fit_transform(df[['Numeric_Column']])

# Display cleaned dataset
print(df.head())
```

Output:

```
===== RESTART: D:/BDA Practical File/Practical 3.py =====
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7 entries, 0 to 6
Data columns (total 5 columns):
#   Column          Non-Null Count  Dtype
---  -
0   ID              7 non-null     int64
1   Name            7 non-null     object
2   Age             7 non-null     int64
3   Salary          6 non-null     float64
4   Numeric_Column  6 non-null     float64
dtypes: float64(2), int64(2), object(1)
memory usage: 412.0+ bytes
None
```

	ID	Name	Age	Salary	Numeric_Column	Normalized_Column
0	1	Alice	25	50000.000000	200.0	0.333333
1	2	Bob	30	73333.333333	150.0	0.000000
2	3	Charlie	35	70000.000000	300.0	1.000000
3	4	David	30	60000.000000	250.0	0.666667
4	5	Eve	35	80000.000000	300.0	1.000000